

Document Title:

From Recursive Production to Governable Method: A Comprehensive Account of the PHAROS Research Program

Date: March 27, 2026

Author: Martin Lepage, PHAROS | pharos-ai.ca

1. Abstract

This document provides a comprehensive account of a research program that began in late 2024 and culminated in the formalization of the PHAROS AI Governance Protocol. The program was driven by a singular scientific intention: to understand how authority is produced, stabilized, and contested within mediated systems—be they ritual, narrative, institutional, or computational. When this inquiry was applied to the researcher's own AI-assisted production process, it encountered a critical failure condition: recursive AI-assisted workflows could generate output faster than they could generate decision clarity, leading to evidence drift, authority inflation, and structural fragility. This document narrates the evolution of the response to that failure. It traces the arc from theoretical and ethnographic beginnings in queer Paganism and digital culture, through the development of a recursive governance architecture, to the rigorous, self-reflexive redesign of that architecture into a deployable, auditable governance method for AI systems. The final output is the PHAROS Recursive AI Governance Protocol, a transdisciplinary science of mediated authority designed to make legitimacy traceable and contestable.

2. The Arc of Inquiry: Research Timeline and Context

The research timeline is best understood as four interconnected phases, each building on the prior one, rather than a simple chronological sequence. The underlying rationale remained stable: to study how claims to truth, legitimacy, and authenticity are assembled, naturalized, and challenged.

Phase 1: Foundational Inquiry (September – December 2024)

This phase established the core objects of study. The work was centered on queer Paganism, neo-shamanism, Wicca, and contemporary ritual life, particularly in Montreal. The rationale was to show how these alternative communities function as "laboratories of legitimacy," revealing how truth claims about identity, lineage, gender, and belonging are produced, mediated by books, commodities, and platforms. The scientific function was epistemic and diagnostic. Parallel work began in cultural and media analysis, using texts like *Buffy the Vampire Slayer* and *Jurassic Park* as "analytic laboratories" to study how power, embodiment, and authority are staged and governed in compressed narrative form. This period was foundational, clarifying the conceptual tools—authenticity, glitch, recursion—that would later be extended to AI systems.

Phase 2: Methodological Elaboration and Laboratory Formation (Spring – Summer 2025)

This phase saw a crucial shift. The research became explicitly **methodological**. The work expanded into two new fronts. First, **public health research** (on DoxyPEP, STI prevention, etc.) brought the question of legitimacy into a more applied institutional setting, revealing how evidence and care structures classify and govern vulnerable populations. Second, the inquiry turned back on itself. Recognizing that its own AI-assisted production was itself a mediated system of authority, the researcher began building controlled production environments—**Hexa, Wheels of Will**—to observe, stress, and govern recursive AI behavior. This was the formation of a "laboratory." The first governance constraints (AI Gov Constitution, Voice Operator) were imposed, marking the moment when the system began to move from pure production to governed production. The work started generating **loop papers**—texts excluded from the main workflow that would later return as diagnostic objects.

Phase 3: Architecture Formalization and Stress Testing (Late Summer – December 2025)

By this stage, the work had moved beyond analyzing mediated systems to using them as part of the inquiry itself. The central problem became **governability under recursion**. The research output was split into distinct classes: **root-method** texts (like *Recursive AI Governance as Executable Method: The Very Long Narrative*) that documented the architecture's construction, and **analytical apex** texts (like *Discursive Authority Under Recursive Pressure*) that theorized the highest-order implications. The architecture was layered into first, second, and third-order controls to manage instability, executable governance, and continuity. The **protocol stack** (Sealed Card, Violet Gem, ALAMBIC, Master Key) was developed to standardize transitions and transformations. This was the **condensation phase**, where governance was formalized from a conceptual orientation into explicit, inspectable, and revisable layers.

Phase 4: Protocol Redesign and Implementation (January – March 2026)

This phase was marked by a rigorous, recursive self-audit. The initial protocol was subjected to a five-pass redesign, each pass diagnosing defects, extracting invariants, and hardening the method against adversarial attacks. The goal was to break the closed evidentiary loop where the method validated itself. The result was the **PHAROS Recursive AI Governance Protocol (v5.0)**, a terminal, stable version with explicit target constructs, evidence hierarchies, provenance rules, and operator controls. Parallel to this, the findings were translated into an operational, public-facing form: the **PHAROS AI Governance Practice**, codified in its **Master Standard Operating Procedure (SOP)**. This phase represents the transition from method-building to institutional implementation, with PHAROS acting as the public-facing surface and CompassAI and Aurora as standalone applications embodying the research.

3. The Method Itself: The PHAROS Recursive AI Governance Protocol

The PHAROS method is not a theory or a policy document; it is an **executable governance architecture**. Its core innovation is to treat governance not as an external overlay on AI production, but as an embedded, recursive part of the production process itself.

The method is structured as a **ten-stage pipeline** that transforms a bounded corpus of materials into interface-bound operational controls. Its key components are:

- **Staged Pipeline:** A linear sequence of stages (Corpus Formation → Intent Embedding → Multi-Model Rendering → Triangulation → Recursive Return → Control Extraction → Determinism Hardening → Deployment Binding → Logging & Audit Trail → Failure Harvesting). Each stage has a declared purpose, a governance question, required deliverables, and a completion condition.
- **Target Constructs:** The protocol measures five core constructs:
 1. **Inferential Carry-Through (TC-1):** The preservation of logical dependency chains across turns.
 2. **Revision Fidelity (TC-2):** Whether a rewrite changes what the passage *implies*, not just its phrasing.
 3. **Perturbation Robustness (TC-3):** The ability to handle a topic interruption and recover the prior thread through structural reference.
 4. **Terminal Path-Dependence (TC-5):** Whether the final conclusion could only have emerged from the specific session trajectory.
- **Evidence Hierarchy:** The protocol establishes a clear hierarchy of evidence, from Level 1 (Direct structural evidence) to Level 4 (Stylistic/lexical noise, which is explicitly inadmissible). This prevents eloquent prose from being mistaken for rigorous reasoning.
- **Control Format:** Any operational control must name a risk pattern, a workflow interface, an accountable actor, an evidence artifact, and a review interval. A control missing any of these is a "governance aspiration."
- **Deterministic Rollup Logic:** The protocol uses 12 deterministic rules to assign an overall status (e.g., ``ready_full``, ``ready_with_bounded_gaps``, ``blocked_not_ready``) based on pass/fail/bounded results from each phase, removing evaluator judgment from the final classification.
- **Anti-Gaming & Anti-Contamination Controls:** The method is explicitly designed to be resistant to "performative compliance." It includes mechanisms like independent spot-checks, verifiable timestamps for probe classifications, and a prohibition on retrospective rationalization of failures.

4. Conceptual and Theoretical Background

The PHAROS method is rooted in a transdisciplinary theoretical framework:

- **Queer Theory & Ritual Studies:** The initial insight that "ritual worlds produce truth claims" is the foundation for analyzing how *any* system (including AI) produces legitimacy. Concepts like **glitch, stutter, and interruption**—derived from queer and performance theory—become diagnostic tools for exposing the seams in systems that claim to be seamless.

- **Digital Culture & Media Analysis:** The method treats media and platforms not as neutral channels but as "infrastructures of mediation" that actively transform the knowledge, identity, and authority they carry. The "laboratory" approach to analyzing texts like *Buffy* and *Jurassic Park* provided a model for treating AI systems as similar sites of compressed social theory.
- **Science & Technology Studies (STS) & Actor-Network Theory (ANT):** The focus on "recursion," "infrastructure," and "distributed agency" (tools that "name through me") is a direct inheritance from STS. The method views governance as an assemblage of humans, non-humans (AI), procedures, and texts.
- **Epistemic Justice & Critical Disability Studies:** The inquiry into AI and authorship is explicitly framed around the question of who gets recognized as a "valid knower." The method seeks to separate rigor from ableist performance norms, treating the body (stutter, illness) not as a contaminant but as a site where the cost of legitimacy is paid.
- **Governance & Institutional Theory:** The final move is a translation of these critical insights into the language of institutional control: risk classification, audit trails, and accountability mechanisms. This is not a retreat from critique but its operationalization.

5. Analysis and Findings

The program's most significant finding is that **recursive AI production is structurally fragile without embedded governance**. The key failures identified are not technical glitches but governance failures:

- **Evidence Drift:** Generated artifacts are treated as source material, eroding the evidentiary basis of future conclusions.
- **Authority Inflation:** Fluency and repetition create an illusion of credibility, leading to "performative compliance" where the form of governance is satisfied without its substance.
- **Method Lock:** Governance procedures become self-referential and stop adapting to new evidence.

The rigorous self-audit of the PHAROS method itself revealed a critical insight: a method that governs recursive AI production must itself be governed recursively. This led to the creation of the **five-pass redesign**, which serves as a case study in how to break a closed evidentiary loop. The process of analyzing the protocol's own defects, applying adversarial stress tests, and patching its vulnerabilities demonstrated that a governance system can only be trustworthy if its own construction is auditable and its limitations are disclosed as residual risks.

The final architecture is thus not a static set of rules but a **"bounded non-final" system**. It expects and accommodates revision, but it structures that revision with explicit gates, provenance, and justifications to prevent uncontrolled drift.

6. Implications for AI Governance at Large

The PHAROS research program offers several profound implications for the broader field of AI governance:

1. **Governance Must Be Recursive:** AI systems are not static; they are part of dynamic, recursive workflows. Governance cannot be a one-time compliance check. It must be a continuous, embedded process that evolves with the system it governs.
2. **Distinguish Form from Substance:** The current field is rife with "performative compliance"—the creation of ethics statements and policies that satisfy a structural requirement but have no operational effect. PHAROS demonstrates that governance must be operationalized into interface-bound controls with named actors, evidence artifacts, and review intervals. A control is not a policy; it is a procedure.
3. **Treat Lineage and Provenance as First-Order Concerns:** The single most important artifact in any governance-sensitive system is the audit trail. PHAROS's focus on source/derived distinctions, hash verification, and deterministic identifiers shows how to make AI workflows reconstructable. Without traceability, trust is impossible.
4. **Incorporate Failure as Evidence:** A robust governance system learns from its own gaps. PHAROS's "failure doctrine" treats abandoned branches, uneven oversight, and archive instability as governance evidence, not noise. This is a direct counter to the industry's tendency to suppress or hide failures.
5. **Acknowledge and Disclose Residual Risks:** No governance system can be perfect. The PHAROS protocol explicitly names its residual risks (e.g., the TC-1 archive confound, template compliance gaming) and provides monitoring approaches. This honesty is more valuable than a false claim of completeness.
6. **Translate Critical Theory into Operational Language:** The program demonstrates that concepts from queer theory, STS, and critical media studies are not merely academic. They can be operationalized into executable governance tools to make power traceable and contestable. The line from "ritual worlds produce truth claims" to an enforceable control register is a direct, albeit long, one.

7. Conclusion

This research program began by studying how authority is made in marginalized ritual spaces. It ended by building a method to govern authority in one of the most powerful technological systems of our time. The arc from **queer Paganism to AI governance** is not a contradiction; it is a coherent scientific trajectory. It is grounded in the conviction that to understand power, one must study it where it is most visible (on the margins) and most urgent (in the infrastructures that govern modern life). The PHAROS method is the concrete output of this conviction: a rigorous, self-reflexive, and executable science of mediated authority, designed to ensure that the systems we build remain accountable to the people they serve.