

AI SOURCE INTEGRITY FIELD GUIDE

Because “the AI said so” is not a citation.

PHAROS AI — v1.0

Validated in the LegiPro Mirofish Evidence Pilot (34 dossiers, 68 turns)

Elemental Agents multi-agent validation architecture

A compact reference guide that teaches researchers and consultants how to trace, verify, and properly attribute AI-assisted claims across seven source categories — with classification matrix, citation repair worksheet, 120+ red-flag glossary, and decision tree.

Core principle: Every substantive claim must carry a reconstructible provenance packet. When that packet cannot be assembled to “Verified” standard, the claim must be labeled, scoped, or rejected before it enters the record.

Executive Summary

Academic integrity policies, journal requirements, funders, and regulators (Law 25, emerging federal AI rules, ISO 42001) are tightening expectations around AI-use disclosure and source accountability. Style guides now tell you *how* to cite AI. This guide tells you *whether* you should cite it at all — and how to trace, verify, classify, and repair claims when the provenance is imperfect. The guide operationalizes seven source categories, a classification matrix, citation repair protocols, a 120+ entry red-flag glossary (grouped for rapid lookup), and a decision tree for when to treat model output as working note versus primary evidence.

Pilot grounding: In the 68-turn replay, debate was required on 38 turns. The most common failure modes were citation drop, evidence drift (claim exceeding source scope), authority inflation, and thin coverage presented as settled. This guide turns those observed failure modes into daily practitioner controls.

1. The Seven Source Categories

Every AI-assisted claim ultimately rests on one or more of these provenance layers. The category determines the verification burden and the allowable claim strength.

Cat	Category	Max Safe Claim	Risk if Unverified
1	Primary peer-reviewed / journal / monograph	Verified (direct) or narrow Inferred	High — many fabricated citations
2	Official statutes, regulations, gov reports, standards	Verified for direct; Inferred for standard extensions	Medium-High (volatility)
3	Corporate primary records (filings, registries, FEC)	Verified for numeric facts & direct statements	Medium
4	Court / tribunal / administrative decisions	Verified for holding; narrow Inferred for ratio	Medium-High
5	News, press releases, trade publications (secondary)	Plausible for event facts only	High
6	Preprints, working papers, unreviewed datasets	Plausible at best; heavy qualification	Very High
7	Model-internal knowledge / uncited synthesis	Never Verified. Often Hallucinated.	Critical

Matrix rule of thumb: If you cannot name the exact document + locator + retrieval date for a Category 1–4 claim, downgrade at least one tier.

2. Source Classification & Verification Matrix (Core Steps)

- 1. Locate claimed source via independent search (DOI/URL + open PDF or official text).
 - 2. Match exact language or data (verbatim quote or number + page/para).
 - 3. Confirm date & version (publication predates model output; current version used).
 - 4. Check for corrections/retractions/amendments (Retraction Watch + agency site).
 - 5. Cross-check independent route (second database, library catalog, or official registry).
 - 6. Human sign-off: named reviewer with domain authority confirms label.
- Label outcome: Verified (direct match) / Inferred (logical extension within bounds) / Plausible (consistent but weak) / Hallucinated (no support or contradiction).

3. Citation Repair Worksheet (Common Failures → Action)

Fabricated or non-resolving DOI/URL: Search exact title+author+year. Nothing matches → Hallucinated. Similar but different doc → Plausible (drift). Rewrite or reject.

Quote does not appear / altered: Search PDF. Absent → Hallucinated. Nearby paraphrase → Inferred or Plausible.

Claim adds scope/numbers/causation (evidence drift): Isolate added element. Necessary from source + uncontroversial premise → Inferred (label). Otherwise Plausible.

Authority inflation: Replace with actual source strength. No specific study/holding → Plausible or Hallucinated.

Secondary presented as primary: Trace footnote to primary. Only secondary exists → Inferred at best.

Model cutoff / post-cutoff 'source': Claimed pub date after cutoff or generation → Hallucinated.

4. Red-Flag Phrase Glossary (Practitioner Edition — 120+ entries)

Grouped for speed. Surface signals validated in the LegiPro pilot (strong correlation with debate-required turns and reclassifications).

Certainty & Consensus Inflat

experts agree • the consensus is • researchers widely hold • it is well established that • scholarship demonstrates • the literature is clear • courts have uniformly • authorities are in agreement • it is beyond dispute • no serious scholar disputes • the position is settled • widely accepted among practitioners

Scope & Generalization Creep

all • every • always • never • in all cases • without exception • any reasonable • invariably • categorically • as a rule • in practice this means • the rule requires that you • one must always

Attribution Fudges

according to sources • reports indicate • it has been reported • available data suggest • studies have shown • evidence indicates • official figures show • internal documents reveal • a review of the literature • a growing body of research

Hedge Disappearance

source 'may/might/could/suggests/appears to/is consistent with/preliminary/limited evidence' → output 'will/does/is/requires/establishes/proves/confirms/settled'

Temporal / Cutoff Violations

as of 2026 • in the latest guidance (2026) • recently amended (post-2025) • the 2026 reform provides • effective Q3 2026 • the platform went live on [future date] • updated circular of March 2026

Sycophancy / Client-Pressure Framing

as you correctly noted • building on your point • to align with your objective • this supports the approach you outlined • the data confirms your hypothesis • this will strengthen your position with the client • as anticipated in your instructions

False Precision

exact percentages/euro amounts/dates/counts presented without source locator when underlying source uses ranges or 'approximately'

Provenance Laundering

as previously established • as discussed above • per our earlier analysis • consistent with the file • as we have seen • building on the prior turn

Legal / Regulatory Citation Errors (TVA, e-invoicing, rescrit, PCG, BTP, FEC, PDP)

article X of the CGI provides that... • per BOFiP §Y... • the DGFIP circular of [wrong date] states... • the PDP registry requires... • under the new Factur-X rules... • the rescrit response confirms... • SIRET must be... • the FEC import validates... • the expert-comptable attestation is... • per the 2025 reform... • the CA3 line requires... • intracom VAT number is sufficient for...

Numerical & Quantitative Hallucinations

the rate is exactly 5.5% • turnover exceeded €2.3m • the fine is 0.5% per day • compliance rate reached 87% • the register contains 14,392 entries • the average is exactly... • median of X • growth of Y%

Entity, Identifier & Registry Fabrication

invented SIREN/SIRET, VAT numbers, PDP identifiers, company names, director names, addresses that do not resolve • the company holds PDP accreditation number... • registered under SIRET 123 456 789 00012 • non-existent circulars, BOFiP references, or journal articles with plausible-sounding titles

Translation & Cross-Lingual Artifacts

awkward literal translations of French admin phrasing presented as natural English • 'the administration considers that...', 'it results from the texts that...' • false friends: rescrit, attestation, ventilation • flattened nuance or reversed meaning

Consensus, Survey & Literature Overclaims

a review of the literature shows... • stakeholder consultations indicate... • multiple studies confirm... • empirical evidence demonstrates... • surveys of practitioners reveal... • the position is widely held among expert-comptables

Causal, Normative & 'Requires' Overreach

this causes... • the rule requires that you... • failure to do X will result in... • the only compliant approach is... • best practice dictates... • to avoid penalties one must...

False or Weak Attribution to Authorities

the Cour des comptes has noted... • the Conseil d'État has ruled... • the OEC recommends... • per the CNIL position... • the AMF has clarified...

Statistical & Data Presentation Tells

precise-looking tables or figures generated without underlying dataset citation or methodology • the sample of N=... with no description of sampling frame, response rate, or limitations

Cross-Border & Multi-Jurisdictional Drift

applying French TVA/PCG rules to Belgian, Swiss, or EU-wide contexts without source for the transposition • 'harmonized under EU rules...' when the specific directive or member-state implementation differs materially

Full designed PDF edition contains the complete curated 120+ entries with usage examples, misapplication warnings, and cross-references to the 7 categories and 4-tier rubric.

5. Decision Tree: AI Output vs. Primary Source

Start with any substantive claim in AI output

↓

Can I retrieve a primary source (Cat 1-4) that directly contains the claim language or data?

YES → Does the retrieved text match verbatim or with only minor, disclosed paraphrase?

YES → **VERIFIED** (use with citation)

NO → Does the claim add interpretation/scope/numbers not present?

YES → **PLAUSIBLE** (flag + narrow)

NO → **INFERRED** (label as "inferred from [source]")

NO → Is the claim a necessary logical consequence of the retrieved primary + one uncontroversial premise?

YES → **INFERRED** (label it)

NO → Does the claim stay within the outer bounds of what a reasonable reader would treat as supported?

YES → **PLAUSIBLE** (flag + provenance)

NO → **HALLUCINATED** (reject; trace origin)

Rule: If you reach "Plausible" or "Hallucinated", do not let the claim enter the final deliverable without explicit labeling and, for client or regulatory work, a human sign-off that accepts the residual risk.

6. Building the Provenance Packet (Minimum for Non-Verified Claims)

1. Exact claim text as it appears in the output.
2. The epistemic label (Inferred / Plausible / Hallucinated) + one-sentence justification.
3. Source excerpt(s) — screenshot or copy of the relevant paragraph + bibliographic header.
4. Retrieval metadata: database/tool used, date/time of retrieval, exact search string or DOI.
5. Reviewer name + date + any override or additional verification performed.
6. (Optional but recommended) Prompt excerpt and model/version if the claim originated in a generative turn.

The LegiPro pilot produced candidate_rollup.md and per-turn JSONL precisely to make this packet reconstructible. Your daily practice should produce the same artifact at claim granularity.

7. Worked Examples (Condensed from Pilot-Style Patterns)

Example A — Thin coverage turned drift (TVA / BTP style)

Output: "The mixed-rate attestation for BTP works must now be filed via the PDP platform under the 2025 e-invoicing rules."

Retrieved source: DGFIP guidance that "envisages" phased adoption for certain BTP operations; no universal "must". Classification: **Plausible (scope overclaim)**. Repair: "DGFIP guidance of [date] contemplates PDP routing for certain BTP mixed-rate operations; exact scope and timetable remain subject to forthcoming implementing measures." Label: Inferred from guidance + explicit limitation.

Example B — Authority inflation on rescrit

Output: "The rescrit procedure does not apply to fully automated fiscal decisions; this is settled administrative practice."

Actual sources: One prior turn noted "current administrative practice in the sampled bureaux has been to decline..."; no statute or published decision stated a general prohibition. Classification: **Plausible** → **Hallucinated** on the "settled" and universal claim. Repair: "In the dossiers reviewed, administrative practice has been to decline rescrit requests concerning fully automated rules; no general statutory bar was located."

(Additional examples in the full PDF edition with before/after text and exact provenance packet entries from the LegiPro replay.)

8. Quick-Reference Card (Printable — Top Daily Checks)

1. [] Exact source title/author/year/journal or statute/article located via independent search.
2. [] DOI/URL/identifier resolves to the claimed document and it is the primary (not a summary).
3. [] Direct quote or number matches verbatim in the source at the cited location.
4. [] Claim does not add scope, causation, universality, or numbers absent from the source.
5. [] Source publication date predates the AI generation; no post-cutoff hallucination.
6. No authority-inflation language ("experts agree", "establishes", "requires") beyond the source's own hedging.

7. [] At least one independent cross-check performed (different database, second primary, or adversarial lens).

8. [] For any non-Verified claim: provenance packet (excerpt + label + retrieval metadata + reviewer) is attached or linked.

If any box is unchecked, downgrade the claim and document the gap.

9. Regulatory & Integrity Mapping

Law 25 (Québec): s. 12.1 and s. 22.1 obligations on automated decision-making and AI use require transparency and accountability. A source integrity failure in a disclosed automated process is direct compliance exposure.

ISO 42001 / NIST AI RMF: Documentation of 'data provenance' and 'output verification' are control expectations. This guide supplies the claim-level granularity those frameworks assume but do not detail.

Journal / funder / thesis requirements: Increasingly demand AI-use statements and reproducible methods. A labeled provenance packet satisfies 'methods' and 'limitations' sections.

Professional liability: The pilot's professional-liability rubric (anti-sycophancy, artifact integrity, boundary, citation, route, turn-specificity) maps directly onto the categories and repair steps above.

10. Integration With Companion Tools

- **AI Evidence Audit Checklist (Product 1):** Run the 50-item checklist first. This Field Guide supplies the deeper source-category and citation-repair layer when the checklist flags a non-Verified claim.

- **Hallucination Risk Scorecard (Product 3):** After classification, score the five dimensions on the claim. A 'Plausible' or 'Hallucinated' label should rarely produce a Low total.

- **Prompt Library & Literature Review Protocol (Products 4 & 7):** Use prompts that explicitly request source locators and force separation of 'retrieved' from 'synthesized.'

- **Governance Starter Kit & Policy Template (Products 5 & 8):** Embed the requirement for provenance packets and source-category classification into organizational AI-use policy.

PHAROS AI — Elemental Agents Framework. This guide turns the discipline demonstrated in the LegiPro pilot into a repeatable daily practice. Public provenance: github.com/martinlepage26-bit/pharos-suite (docs/evidence/2026-06-05-legipro-mirofish/). Lifetime updates via Gumroad purchase.